



Fall 2017

CSCI 402

Warmup Assignment #2

(100 points total)

Multi-threading - Token Bucket Emulation in C

Due 11:45PM 9/29/2017 (firm)

This spec is **private** (i.e., only for students who took or are taking CSCI 402 at USC). You do **not** have permissions to **display this spec** at a public place (such as a public bitbucket/github). You also do **not** have permissions to **display the code** you write to implement this spec at a public place since your code was written to implement a private spec. (If a prospective employer asks you to post your code, please tell them that you do not have permissions to do so; but you can send them a **private copy**.)

Assignment Description

(Please check out the [Warmup 2 FAQ](#) before sending your questions to the TAs, the course producers, or the instructor.)

In this assignment, you will emulate/simulate a **traffic shaper** who transmits packets controlled by a **token bucket filter** depicted below using multi-threading within a single process. If you are not familiar with pthreads, you should read Chapter 2 of our [required textbook](#).

IMPORTANT: Please note that this assignment is posted before all the background materials (e.g., **Unix signals**) have been covered in lectures. If you do **not** want to learn about these components on your own (by learning from the textbook), please delay starting this project until they are covered in class. There will be plenty of time to implement this project after the relevant topics are covered in class.




Figure 1: A system with a token bucket filter.

Figure 1 above depicts the **system** you are required to emulate. The **token bucket** has a capacity (bucket depth) of B tokens. Tokens arrive into the token bucket according to an unusual arrival process where the inter-arrival time between two consecutive tokens is $1/r$. We will call r the **token arrival rate** (although technically speaking, it's not exactly the token arrival rate; please understand that this is **quite different** from saying that the tokens arrive at a constant rate of r). Extra tokens (overflow) would simply disappear if the token bucket is full. A token bucket, together with its control mechanism, is referred to as a **token bucket filter**.

Packets arrive at the token bucket filter according to an unusual arrival process where the inter-arrival time between two consecutive packets is $1/\lambda$. We will call λ the **packet arrival rate** (although technically speaking, it's not exactly the packet arrival rate; please understand that this is **quite different** from saying that the packets arrive at a constant rate of λ). Each packet requires P tokens in order for it to be eligible for transmission. (Packets that are eligible for transmission are queued at the $Q2$ facility.) When a packet arrives, if $Q1$ is not empty, it will just get queued onto the $Q1$ facility. Otherwise, it will check if the token bucket has P or more tokens in it. If the token bucket has P or more tokens in it, P tokens will be removed from the token bucket and the packet will join the $Q2$ facility (technically speaking, you are **required** to first add the packet to $Q1$ and timestamp the packet, remove the P tokens from the token bucket and the packet from $Q1$ and timestamp the packet, before moving the packet into $Q2$), and **wake up** the servers in case they are sleeping. If the token bucket does not have enough tokens, the packet gets queued into the $Q1$ facility. (Please note that, **in this case**, you do **not** have to check if there is enough tokens in the bucket so you can move the packet at the head of $Q1$ into $Q2$ and you need to understand why you do **not** need to perform such a check.) Finally, if the number of tokens required by a packet is larger than the bucket depth, the packet must be **dropped** (otherwise, it will block all other packets that follow it).

The transmission facility (denoted as $S1$ and $S2$ in the above figure and they are referred to as the "**servers**") serves packets in $Q2$ in the first-come-first-served order and at a service rate of μ per second. When a server becomes available, it will dequeue the first packet from $Q2$ and start transmitting the packet. When a packet has received $1/\mu$ seconds of service, it leaves the system. You are required to keep the servers **as busy as possible**.

When a token arrives at the **token bucket**, it will add a token into the **token bucket**. If the bucket is already full, the token will be lost. It will then check to see if $Q1$ is empty. If $Q1$ is not empty, it will see if there is enough tokens to make the packet at the head of $Q1$ be eligible for transmission (packets in $Q1$ in also served in the first-come-first-served order). If it does, it will remove the corresponding number of tokens from the token bucket, remove that packet from $Q1$ and move it into $Q2$, and **wake up** the servers in case they are sleeping. It will then check the packet that is now at the head of $Q1$ to see if it's also eligible for transmission, and so on.

Technically speaking, the "servers" are not part of the "token bucket filter". Nevertheless, it's part of this assignment to emulate the servers because the servers are considered part of the "system" to be emulated.

Our system can run in only one of two modes.

Deterministic : In this mode, all inter-arrival times are equal to $1/\lambda$ seconds, all packets require exactly P tokens, and all service times are equal to $1/\mu$ seconds (**all rounded to the nearest millisecond**). If $1/\lambda$ is greater than 10 seconds, please use an inter-arrival time of 10 seconds. If $1/\mu$ is greater than 10 seconds, please use a service time of 10 seconds.

Trace-driven : In this mode, we will drive the emulation using a trace specification file. Each line in the trace file specifies the inter-arrival time of a packet, the **number of tokens** it need in order for it to be eligible for transmission, and its **service time**. (Please note that in this mode, it's perfectly fine if an inter-arrival time or a service time is greater than 10 seconds.)

Your job is to emulate the packet and token arrivals, the operation of the token bucket filter, the first-come-first-served queues $Q1$ and $Q2$, and servers $S1$ and $S2$. You also must produce a trace of your emulation for every important event occurred in your emulation. Please see more details below for the requirements.

You must use:

- one thread for packet arrival
- one thread for token arrival
- one thread for each server

You must **not** use one thread for each packet.

In addition, you must use at least one mutex to protect **Q1**, **Q2**, and the token bucket. (It is recommended that you use **exactly one** mutex to protect **Q1**, **Q2**, and the token bucket.)

Finally, **Q1** and **Q2** must have infinite capacity (i.e., you should use **My420List** from [warmup assignment #1](#) to implement them and **not** use arrays).

We will **not** go over the [slides for this assignment](#) in class. Although it's important that you are familiar with it. Please read it over. If you have questions, please e-mail the **instructor**.

Compiling

Please use a Makefile so that when the grader simply enters:

```
make warmup2
```

an executable named **warmup2** is created (minor variation is permitted if you document it). Please make sure that your submission conforms to [other general compilation requirements](#) and [README requirements](#).

Commandline

The command line syntax (also known as "usage information") for **warmup2** is as follows:

```
warmup2 [-lambda lambda] [-mu mu] [-r r] [-B B] [-P P] [-n num] [-t tsfile]
```

Square bracketed items are optional. You must follow the UNIX convention that **commandline options** can come in any order. (Note: a **commandline option** is a commandline argument that begins with a - character in a commandline syntax specification.) Unless otherwise specified, output of your program must go to stdout and error messages must go to stderr.

The `lambda`, `mu`, `r`, `B`, and `P` parameters all have obvious meanings (according to the description above). The `-n` option specifies the total number of packets to arrive. If the `-t` option is specified, `tsfile` is a [trace specification file](#) that you should use to drive your emulation. In this case, you should ignore the `-lambda`, `-mu`, `-P`, and `-n` commandline options and run your emulation in the [trace-driven mode](#). You may assume that `tsfile` conforms to the [tracefile format specification](#). (This means that if you detect an error in this file, you may simply print an error message and call `exit()`. There is no need to perform error recovery.) If the `-t` option is not used, you should run your emulation in the [deterministic mode](#).

The default value (i.e., if it's not specified in a commandline option) for `lambda` is 1 (packets per second), the default value for `mu` is 0.35 (packets per second), the default value for `r` is 1.5 (tokens per second), the default value for `B` is 10 (tokens), the default value for `P` is 3 (tokens), and the default value for `num` is 20 (packets). `B`, `P`, and `num` must be positive integers with a maximum value of 2147483647 (0x7fffffff). `lambda`, `mu`, and `r` must be positive real numbers.

If $1/r$ is greater than 10 seconds, please use an inter-token-arrival time of 10 seconds.

Running Your Code and Program Output

The emulation should go as follows. At emulation time 0, all 4 threads (arrival thread, token depositing thread, and servers **S1** and **S2** threads) got started. The arrival thread would sleep so that it can wake up at a time such that the inter-arrival time of the first packet would match the specification (either according to `lambda` or the first record in a tracefile). At the same time, the token depositing thread would sleep so that it can wake up every $1/r$ seconds and would try to deposit one token into the token bucket. The actual arrival time of the first packet `p1` is denoted as time `T1`, the actual arrival time of the 2nd packet `p2` is denoted as time `T2`, and so on.

As a packet or a token arrives, or as a server becomes free, you need to follow the [operational rules of the token bucket filter](#). Since we have four threads accessing shared data structures, you must use the tricks you learned from Chapter 2 related lectures. Please also check out the [slides for this assignment](#) for the skeleton code for these threads.

You are required to produce a detailed trace for every important event occurred during the emulation. Each line in the trace must correspond to one of the following situations:

- If a packet is served by a server (server **S1** is assumed below for illustration), there must be exactly 7 output lines that correspond to this packet. They are:

```
p1 arrives, needs 3 tokens, inter-arrival time = 503.112ms
p1 enters Q1
p1 leaves Q1, time in Q1 = 247.810ms, token bucket now has 0 token
p1 enters Q2
p1 leaves Q2, time in Q2 = 0.216ms
p1 begins service at S1, requesting 2850ms of service
p1 departs from S1, service time = 2859.911ms, time in system = 3109.731ms
```

Please note the following:

- The value printed for "inter-arrival time" must equal to the timestamp of the "p1 arrives" event minus the timestamp of the "arrives" event for the previous packet.
 - The value printed for "time in Q1" must equal to the timestamp of the "p1 leaves Q1" event minus the timestamp of the "p1 enters Q1" event.
 - The value printed for "time in Q2" must equal to the timestamp of the "p1 leaves Q2" event minus the timestamp of the "p1 enters Q2" event.
 - The value printed for "requesting ???ms of service" must be the requested service time (which must be an integer) of the corresponding packet.
 - The value printed for "service time" must equal to the timestamp of the "p1 departs from S1" event minus the timestamp of the "p1 begins service at S1" event (and it should be larger than the requested service time printed for the "begin service" event).
 - The value printed for "time in system" must equal to the timestamp of the "p1 departs from S1" event minus the timestamp of the "p1 arrives" event;
- If a packet is dropped, you must print:

```
p1 arrives, needs 3 tokens, inter-arrival time = 503.112ms, dropped
```

Please note that the value printed for "inter-arrival time" must equal to the timestamp of the "p1 arrives" event minus the timestamp of the "arrives" event for the previous packet.

- If <Cntrl+C> is pressed by the user, you must print the following (and print a '\n' before it to make sure that it lines up with all the other trace printouts):

```
SIGINT caught, no new packets or tokens will be allowed
```

Please understand that in order for the above to get printed correctly in a trace printout, using a signal handler to catch signals may not work. You are strongly advised to use a separate SIGINT-catching thread and uses sigwait().

- If a packet is removed when it's in Q# (Q1 or Q2) because <Cntrl+C> is pressed by the user, you must print:

```
p1 removed from Q#
```

- If a token is accepted, you must print:

```
token t1 arrives, token bucket now has 1 token
```

- If a token is dropped, you must print:

```
token t1 arrives, dropped
```

- When you are ready to start your emulation, you must print:

```
emulation begins
```

- When you are ready to end your emulation, you must print:

```
emulation ends
```

All the numeric values above are made up. You must replace them with the actual packet number, actual number of tokens required, actual server number, measured inter-arrival time, measured time spent in Q1, actual number of tokens left behind when a packet is moved into Q2, measured time spent in Q2, measured service time, and measured time in the system.

The output **format** of your program **must** satisfy the following requirements.

- You must first print all the emulation parameters. Please see the [sample printout](#) for what the output must look like.
- Whenever a token arrives, you must assign a number to it, and add it to the token bucket. You must then print its arrival time, the fact that it has arrived, and the number of tokens in the token bucket. Please see the [sample printout](#) for what the output must look like.
- Whenever a packet arrives, you must assign a number to it. You must then print its arrival time, the fact that it has arrived, the number of tokens it needs for transmission, and the time between its arrival time and the arrival time of the previous packet. Please see the [sample printout](#) for what the output must look like.

You then must append this packet onto Q1. Afterwards, you must then print the time this packet entered Q1 and the fact that it has entered Q1. Please see the [sample printout](#) for what the output must look like.

Later on, when this packet leaves Q1, it removes the correct number of tokens from the token bucket. You must then print the time this packet leaves Q1, the fact that it has left Q1, the amount of time it spent in Q1, and the number of tokens in the token bucket. Please see the [sample printout](#) for what the output must look like.

You must then append this packet onto Q2. Afterwards, you must then print the time this packet entered Q2 and the fact that it has entered Q2. Please see the [sample printout](#) for what the output must look like.

Later on, when this packet leaves Q2 and enters the server, you must then print which server the packet entered, the time the packet began service, the fact that it has begun service, and the amount of time it spent in Q2. Please see the [sample printout](#) for what the output must look like.

- When emulation ends, you must print all the necessary statistics. Please see the [sample printout](#) for what the output must look like. If a particular statistics is not applicable (e.g., will cause divide-by-zero error), instead of printing a numeric value, please print "N/A" followed by an explanation (such as "no packet arrived at this facility"). Please note that your program output must never contain any "NaN" (which means "not-a-number").

Below is **an example** what your program output must look like (please note that the values used here are just a bunch of unrelated random numbers for illustration purposes):

```
Emulation Parameters:
  number to arrive = 20
  lambda = 2          (print this line only if -t is not specified)
  mu = 0.35           (print this line only if -t is not specified)
  r = 4
  B = 10
  P = 3               (print this line only if -t is not specified)
  tsfile = FILENAME   (print this line only if -t is specified)

00000000.000ms: emulation begins
00000251.726ms: token t1 arrives, token bucket now has 1 token
00000502.031ms: token t2 arrives, token bucket now has 2 tokens
00000503.112ms: p1 arrives, needs 3 tokens, inter-arrival time = 503.112ms
00000503.376ms: p1 enters Q1
00000751.148ms: token t3 arrives, token bucket now has 3 tokens
00000751.186ms: p1 leaves Q1, time in Q1 = 247.810ms, token bucket now has 0 token
00000752.716ms: p1 enters Q2
00000752.932ms: p1 leaves Q2, time in Q2 = 0.216ms
```

```

00000752.982ms: p1 begins service at S1, requesting 2850ms of service
00001004.271ms: p2 arrives, needs 3 tokens, inter-arrival time = 501.159ms
00001004.526ms: p2 enters Q1
00001007.615ms: token t4 arrives, token bucket now has 1 token
00001251.259ms: token t5 arrives, token bucket now has 2 tokens
00001505.986ms: p3 arrives, needs 3 tokens, inter-arrival time = 501.715ms
00001506.713ms: p3 enters Q1
00001507.552ms: token t6 arrives, token bucket now has 3 tokens
00001508.281ms: p2 leaves Q1, time in Q1 = 503.755ms, token bucket now has 0 token
00001508.761ms: p2 enters Q2
00001508.874ms: p2 leaves Q2, time in Q2 = 0.113ms
00001508.895ms: p2 begins service at S2, requesting 1900ms of service
...
00003427.557ms: p2 departs from S2, service time = 1918.662ms, time in system = 2423.286ms
00003612.843ms: p1 departs from S1, service time = 2859.861ms, time in system = 3109.731ms
...
????????.???ms: p20 departs from S?, service time = ????.???ms, time in system = ????.???ms
????????.???ms: emulation ends

```

Statistics:

```

average packet inter-arrival time = <real-value>
average packet service time = <real-value>

average number of packets in Q1 = <real-value>
average number of packets in Q2 = <real-value>
average number of packets at S1 = <real-value>
average number of packets at S2 = <real-value>

average time a packet spent in system = <real-value>
standard deviation for time spent in system = <real-value>

token drop probability = <real-value>
packet drop probability = <real-value>

```

In the Emulation Parameters section, please print the emulation parameters specified by the user or the default values mentioned above. Please do **not** print the "adjusted" values because certain parameters are too small. (For example, if λ is 0.01, you must print 0.01 and not 0.1.)

After Emulation Parameters section comes the Event Trace section. The first column there contains timestamps and they correspond to event times, measured relative to the start of the emulation. You need to figure out how to make sure that the timestamp values look reasonable (e.g., never decrease in value). Please use 8 digits (with leading zeroes) to the left of the decimal point and 3 digits after the decimal point for all the timestamps in this column. All time intervals must be printed in milliseconds with 3 digits after the decimal point. In the printout, after emulation parameters, all values reported must be **measured** values.

In the Statistics section, the **average number of packets** at a facility can be obtained by adding up all the time spent at that facility (for all relevant packets) divided by the total emulation time. The **time spent in system** for a packet is the difference between the time the packet departed from the server and the time that packet arrived. The **token drop probability** is the total number of tokens dropped because the token bucket was full divided by the total number of tokens that was produced by the token depositing thread. The **packet drop probability** is the total number of packets dropped because the number of tokens required is larger than the bucket depth divided by the total number of packets that was produced by the arrival thread.

All real values in the Emulation Parameters and Statistics sections must be printed with at least 6 significant digits. (If you are using `printf()`, you can use `("%.6g")`.) A timestamp in the beginning of a line of trace output must be in milliseconds with 8 digits (zero-padded) before the decimal point and 3 digits (zero-padded) after the decimal point.

Please use **sample means** when you calculated the averages. If n is the number of sample, this mean that you should divide things by n (and not $n-1$).

The unit for time related *statistics* must be in seconds (and not milliseconds).

Let X be something you measure. The standard deviation of X is the square root of the variance of X . The variance of X is the average of the square of X minus the square of the average of X . Let $E(X)$ denote the average of X , you can write:

$$\text{Var}(X) = E(X^2) - [E(X)]^2$$

If the user presses <Ctrl+C> on the keyboard, you must stop the arrival thread and the token depositing thread, remove all packets in **Q1** and **Q2**, let your server finish serving the current packet in the usual way, and output statistics in the usual way. (Please note that it may not be possible to remove all packets in **Q1** at the instance of signal delivery. The idea here is that once signal delivery has occurred, the only packet you should serve are the ones in service. All other packets should be removed from the system.)

You can divide the packets into 3 categories.

1. **Completed packets:** these are the packets that made it all the way to the server and completed service at the server.
2. **Dropped packets:** these are the packets arrived into the system but never made it even to **Q1** because it needs too many tokens.
3. **Removed packets:** these are the packets that got into **Q1** to begin with but never made it to the server.

All packets should participate in the calculation of the average packet inter-arrival time and packet drop probability statistics. Only completed packets should participate in the calculation of the average packet service time statistics. Only completed packets should participate in the calculation of the average number of packets in **Q1/Q2/S1/S2** and time spent in system statistics.

Finally, when no more packet can arrive into the system, you must stop the arrival thread as soon as possible. Also, when **Q1** is empty and no future packet can arrive into **Q1**, you must stop the token depositing thread as soon as possible.

Trace Specification File Format

The trace specification file is an ASCII file containing $n+1$ lines (each line is terminated with a "\n") where n is the total number of packets to arrive. Line 1 of the file contains a positive integer which corresponds to the value of n . Line k of the file contains the inter-arrival time in milliseconds (a positive integer), the number of tokens required (a positive integer), and service time in milliseconds (a positive integer) for packet $k-1$. The 3 fields are separated by space or tab characters (or any combination of any number of these characters). There must be no leading or trailing space or tab characters in a line. If a line is longer than 1,024 characters (including the '\n' at the end of a line), it is considered an error. [A sample tsfile](#) for $n=3$ packets is provided. It's content is listed below:

```
3
2716  2   9253
7721  1  15149
972   3   2614
```

In the above example, packet 1 is to arrive 2716ms after emulation starts, it needs 2 tokens to be eligible for transmission, and its service time should be 9253ms; the inter-arrival time between packet 2 and 1 is to be 7721ms, it needs 1 token to be eligible for transmission, and its service time should be 15149ms; the inter-arrival time between packet 3 and 2 is to be 972ms, it needs 3 token to be eligible for transmission, and its service time should be 2614ms.

In the above example, you should treat these numeric values as "**targets**" or your emulation. In your trace output, you need to print **what you measured** (i.e., by reading the clock). It should be very unlikely that a **measured inter-arrival time** or a **measured service time** has **exactly the same value** as its corresponding target value. For example, the inter-arrival time of packet 3 is suppose to be 972 milliseconds. If the reported actual inter-arrival time between packets 2 and 3 is exactly 972.000 milliseconds, you should look for bugs in your code! Actually, you should probably get a different value every time you rerun your emulation.

This file is expected to be error-free. (This means that if you detect a real error in this file, you must simply print an error message and call `exit()`. There is no need to print statistics or perform error recovery.)

Grading Guidelines

The [grading guidelines](#) has been made available. Please run the scripts in the guidelines on `nunki.usc.edu`. It is possible that there are bugs in the guidelines. If you find bugs, please let the instructor know as soon as possible. (**Note:** the grading guidelines is subject to change without notice.)

Please note that although the grader will follow the grading guidelines to grade, the grader may use a different set of trace files and commandline arguments.

The grading guidelines is the **only** grading procedure we will use to grade your program. No other grading procedure will be used. Please note that the grader may use a **different** set of **trace files** and **commandline arguments**. (We may also make minor changes if we discover bugs in the script or things that we forgot to test.) It is strongly recommended that you run your code through the scripts in the grading guidelines.

Miscellaneous Requirements & Hints

- Please read the [general programming FAQ](#) if you need a refresher on file I/O and bit/byte manipulations in C.
- You must **NOT use any external code segments** to implement this assignment. You must implement all these functionalities from scratch.
- You must **NOT use semaphores** to implement this assignment. You must implement thread synchronization using pthread mutexes and condition variables.
- Please do not use an array to store all the packets. If you do that, you may end up losing a lot of points because there will be cases your program will not be able to handle. Please design your program so that it can handle billions of packets. If you don't know how to do this, you should start a discussion in the class Google Group.
- You are required to use [separate compilation](#) to compile your source code. You must divide your source code into separate source files in a logical way. You also must **not** put the bulk of your code in header files!
- If you have `.nfs*` files you cannot remove, please see [notes on .nfs files](#).
- Please use `gettimeofday()` to get time information with a **microsecond** resolution. You can use `select()` or `usleep()` (or equivalent) to sleep for a specified number of **microseconds**.
- I do **not** recommend using a signal handler to catch SIGINT. You are strongly advised to use a separate SIGINT-catching thread and uses `sigwait()`.
- Your code must not "busy-wait"! "Busy-waiting" means that you have code that's staying in a tight loop to wait for some condition to become true. Such code is very unfriendly to the environment you are running your program in. Therefore, you must not do busy-waiting. If you find a piece of busying waiting code, just insert `"usleep(100000)"` inside the loop to sleep for 0.1 second before checking the condition again. Since it's so easy to not do busy-waiting, if the grader sees that your code is doing busy-waiting, a lot of points will be deducted.

If your code is not doing anything useful (i.e., just waiting for something) and you run "top" from the commandline on `nunki.usc.edu` and you see that your program is taking up one of the top spots in CPU percentages and showing something like more than 0.5%, there is a good chance that you have busy-waiting code. In this is the case, run your code under `gdb` and run "top" in another terminal. When your code is not doing anything useful and your process starts to show up in

"top", press <Ctrl+C> in gdb. Hopefully, your program will break inside your busy-waiting loop (use the where command to see where you are inside gdb)!

- Please understand the meaning of **inter-arrival time**. It means **time between consecutive packets**. Let's say the inter-arrival time is 1 second and we are running in the deterministic mode. Does it mean that packet 1 will arrive at 1 second after the start of emulation, packet 2 will arrive at 2 second after the start of emulation, packet 3 will arrive at 3 second after the start of emulation, etc.? No. You need to **schedule** packet 1 to arrive at 1 second after the start of emulation. But packet 1 will most likely **not** arrive until a little bit later than 1 second after the start of emulation. You record the **actual time of arrival** of packet 1 and you add 1 second to it and that would give you the **expected arrival time** of packet 2. Then you can **schedule** packet 2 to arrive at its **expected arrival time**.

Here are some additional hints:

- For this assignment, you are implementing a **time-driven** emulation (and *not* an event-driven simulation such as the ns-2 in networking). For an event-driven emulation, you can easily implement this project using a single thread and an event queue. In a time-driven emulation, a thread must sleep for the amount of time that it supposes to take to do the a job. For example, if a server would take 317 milliseconds to serve a job, it would actually sleep (using `select/usleep()`) for some period of time so that the packet seem to stay in the server for 317 milliseconds. Similarly, if the arrival thread needs to wait 634 milliseconds between the arrivals of packets p1 and p2, it should sleep for some period of time so that it looks like packet p2 arrives 634 milliseconds after packet p1 has arrived.
- You need to calculate time correctly for the `select/usleep()` call mentioned above. For example, if the arrival thread needs to wait 634 milliseconds between the arrivals of packets p1 and p2. Assuming that it took 45 milliseconds to do bookkeeping and to enqueue the packet to the queueing system, you should sleep for 589 milliseconds (and not sleep for 634 milliseconds). (Please note such calculation does not apply to sleeping for the server thread but only applies to the packet arrival and token depositing threads.)
- Continue with the above example, if it took more than 634 milliseconds to do bookkeeping and to enqueue the packet to the queueing system, what should you do? Since you cannot sleep for a negative number of microseconds, you should skip sleeping. This is "**best-effort**" emulation. You are **not required** to hit the target time, you just need to try your best to match inter-arrival times. When you cannot, you should still try your best (i.e., 0 is the closest number to a negative number).

Submission

All assignments are to be submitted electronically - including your README file. To submit your work, you must first tar all the files you want to submit into a **tarball** and gzip it to create a **gzipped tarfile** named **warmup2.tar.gz**. Then you upload **warmup2.tar.gz** to the [Bistro](#) system. On [nunki.usc.edu](#) or [aludra.usc.edu](#), the command you can use to create a gzipped tarfile is:

```
/usr/usc/bin/gtar cvzf warmup2.tar.gz MYFILES
```

Where **MYFILES** is the list of file names that you are submitting (you can also use wildcard characters if you are sure that it will pick up only the right files). **DO NOT** submit your compiled code, just your source code and README file. **Two point will be deducted** if you submit extra binary files, such as **warmup2**, **.o**, **core**, or files that can be **generated** from the rest of your submission.

Please note that the 2nd commandline argument of the `gtar` command above is the **output** filename of the `gtar` command. So, if you omit **warmup2.tar.gz** above, you may accidentally replace one of your files with the output of the `gtar` command. So, please make sure that the first commandline argument is **cvzf** and the 2nd commandline argument is **warmup2.tar.gz**.

A [w2-README.txt template file is provided here](#). You must save it as your w2-README.txt file and fill it out with your documentation information (i.e., **replace all "(Comments?)"** with your evaluation and **replace all standalone "?"** with information appropriate for your submission).

Here is a sample command for creating your warmup2.tar.gz file (your command will vary depending on what files you want to submit):

```
/usr/usc/bin/gtar cvzf warmup2.tar.gz *.c *.h Makefile w2-README.txt
```

You should read the output of the above commands carefully to make sure that warmup2.tar.gz is created properly. If you don't understand the output of the above commands, you need to learn how to read it! It's your responsibility to ensure that warmup2.tar.gz is created properly.

You need to run [bsubmit](#) to submit warmup2.tar.gz to the submission server. Please use the following command:

```
~csci551b/bin/bsubmit upload \  
-email `whoami`@usc.edu \  
-event merlot.usc.edu_80_1372906710_171 \  
-file warmup2.tar.gz
```

Please note that the quotation marks surrounding whoami are **back-quote** characters and not single quotes. It's best if you just copy and paste the above command into your console and not try to type the whole command in to avoid mistakes.

If the command is executed successfully, the output should look like the [sample mentioned in the submission web page](#). If it doesn't look like that, please fix your command and rerun it until it looks right. If there are problems, please contact the instructor.

It is extreme important that you also [verify your submission](#) after you have submitted warmup2.tar.gz electronically to make sure that everything you have submitted is everything you wanted us to grade.

Finally, please be familiar with the [Electronic Submission Guidelines](#) and information on the [bsubmit web page](#).

[Last updated Mon Aug 07 2017] [Please see [copyright](#) regarding copying.]